

In [76]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import matplotlib.ticker as tcr
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import precision_score, recall_score, accuracy_score, f1_score
from sklearn.metrics import roc_curve, auc
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import validation_curve
```

In [77]:

```
data = pd.read_csv('/home/amybirdee/hobby_projects/propensity_to_buy/Propensity_data.csv', delimiter = ',', \
                    low_memory = False)
```

Preliminary data checks


In [78]:

```
#the data consist of a user id and 0s and 1s to show whether they clicked on various items on the page and ultimately  
#whether or not they placed an order  
data.head()
```

Out[78]:

	UserID	basket_icon_click	basket_add_list	basket_add_detail	sort_by	image_picker	a
0	a720-6b732349-a720-4862-bd21-644732	0	0	0	0	0	
1	b74a-6b737717-b74a-45c3-8c6a-421140	0	1	0	1	0	
2	7775-6b73b976-7775-4324-b1d9-622031	0	0	0	0	0	
3	4b8a-6b74bd36-4b8a-4d10-a008-67143	0	0	0	0	0	
4	7009-6b768104-7009-4526-9da6-129024	0	0	0	0	0	

5 rows × 21 columns



In [79]:

```
#232,000 rows of data  
data.shape
```

Out[79]:

(232239, 21)

In [80]:

```
#don't need the id column so dropping this
data = data.drop('UserID', axis = 1)
data.head()
```

Out[80]:

	basket_icon_click	basket_add_list	basket_add_detail	sort_by	image_picker	account_page_click
0	0	0	0	0	0	0
1	0	1	0	1	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0

In [81]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 232239 entries, 0 to 232238
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   basket_icon_click                    232239 non-null  int64
1   basket_add_list                      232239 non-null  int64
2   basket_add_detail                    232239 non-null  int64
3   sort_by                             232239 non-null  int64
4   image_picker                        232239 non-null  int64
5   account_page_click                  232239 non-null  int64
6   promo_banner_click                  232239 non-null  int64
7   detail_wishlist_add                 232239 non-null  int64
8   list_size_dropdown                  232239 non-null  int64
9   closed_minibasket_click             232239 non-null  int64
10  checked_delivery_detail              232239 non-null  int64
11  checked_returns_detail              232239 non-null  int64
12  sign_in                             232239 non-null  int64
13  saw_checkout                        232239 non-null  int64
14  saw_sizecharts                      232239 non-null  int64
15  saw_delivery                        232239 non-null  int64
16  saw_account_upgrade                 232239 non-null  int64
17  saw_homepage                       232239 non-null  int64
18  returning_user                      232239 non-null  int64
19  ordered                             232239 non-null  int64
dtypes: int64(20)
memory usage: 35.4 MB
```

Exploratory data analysis

In [82]:

```
#checking the proportion of customers who place an order - only 8% place an order so th  
is must be a high ticket item. Will  
#assume this is an interactive exercise bike for the rest of the analysis  
order = data.groupby('ordered').size().to_frame().reset_index().rename(columns = {0: 'c  
ount'})  
order
```

Out[82]:

	ordered	count
0	0	213146
1	1	19093

In [149]:

```
#creating chart
fig = plt.figure(figsize = (15, 6))
ax = plt.subplot(1, 2, 1)

_, _ = plt.pie(order['count'], colors = ['red', 'black'], radius = 1.5, auto
pct = ('%.0f%%'), \
        counterclock = False, startangle = -270)

plt.legend(labels = ['No order', 'Order'], loc = 'right', bbox_to_anchor = (1.5, 0.7),
fontsize = 12)

#setting the colour of percentage labels to white
for autotext in autotexts:
    autotext.set_color('white')
    autotext.set_weight('bold')
    autotext.set_fontsize(12)

plt.title('Proportion of customers who placed an order with Pedal Power', y = 1.2, font
size = 12)

#second chart
ax = plt.subplot(1, 2, 2)

plt.bar(order['ordered'], order['count'], color = 'red')

#removing chart borders
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

#adding data labels for bars
bars = plt.bar(order['ordered'], order['count'], color = 'red')

for bar in bars:
    yval = bar.get_height()
    #the '{:,}' command adds a thousand separator to the labels
    ax.annotate('{:,}'.format(yval),
    xy = (bar.get_x() + bar.get_width() / 2, yval),
    #shows label position on x and y axis
    xytext = (0, 3),
    textcoords = 'offset points', ha = 'center', va = 'bottom', fontsize = 12)


#function to add comma separator to labels
def comma(x, pos):
    return format(x, "6,.0f")

#this code adds a comma separator to the y tick marks
ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))

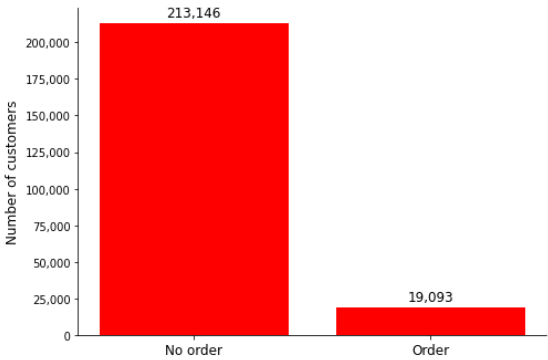
plt.xticks(order['ordered'], labels = ['No order', 'Order'], fontsize = 12)
plt.ylabel('Number of customers', fontsize = 12)
plt.title('Number of customers who placed an order with Pedal Power', y = 1.2, fontsize
= 12)
plt.tight_layout()

plt.savefig('order_proportion_and_volumes')
```

Proportion of customers who placed an order with Pedal Power



Number of customers who placed an order with Pedal Power



In [84]:

```
#grouping by customers who added the bike to their wishlist and/or placed an order
wishlist = data.groupby(['detail_wishlist_add', 'ordered']).size().to_frame().reset_index().\
rename(columns = {0: 'count'})
wishlist
```

Out[84]:

	detail_wishlist_add	ordered	count
0	0	0	211508
1	0	1	18899
2	1	0	1638
3	1	1	194

In [85]:

```
#filtering for users who added to wishlist only
wishlist = wishlist[wishlist['detail_wishlist_add'] == 1].reset_index(drop = True)
wishlist
```

Out[85]:

	detail_wishlist_add	ordered	count
0	1	0	1638
1	1	1	194

In [153]:

```
#creating chart
fig = plt.figure(figsize = (8, 6))
ax = plt.subplot(1, 1, 1)

plt.bar(wishlist['ordered'], wishlist['count'], color = 'red')

#removing chart borders
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

#adding data labels for bars
bars = plt.bar(wishlist['ordered'], wishlist['count'], color = 'red')

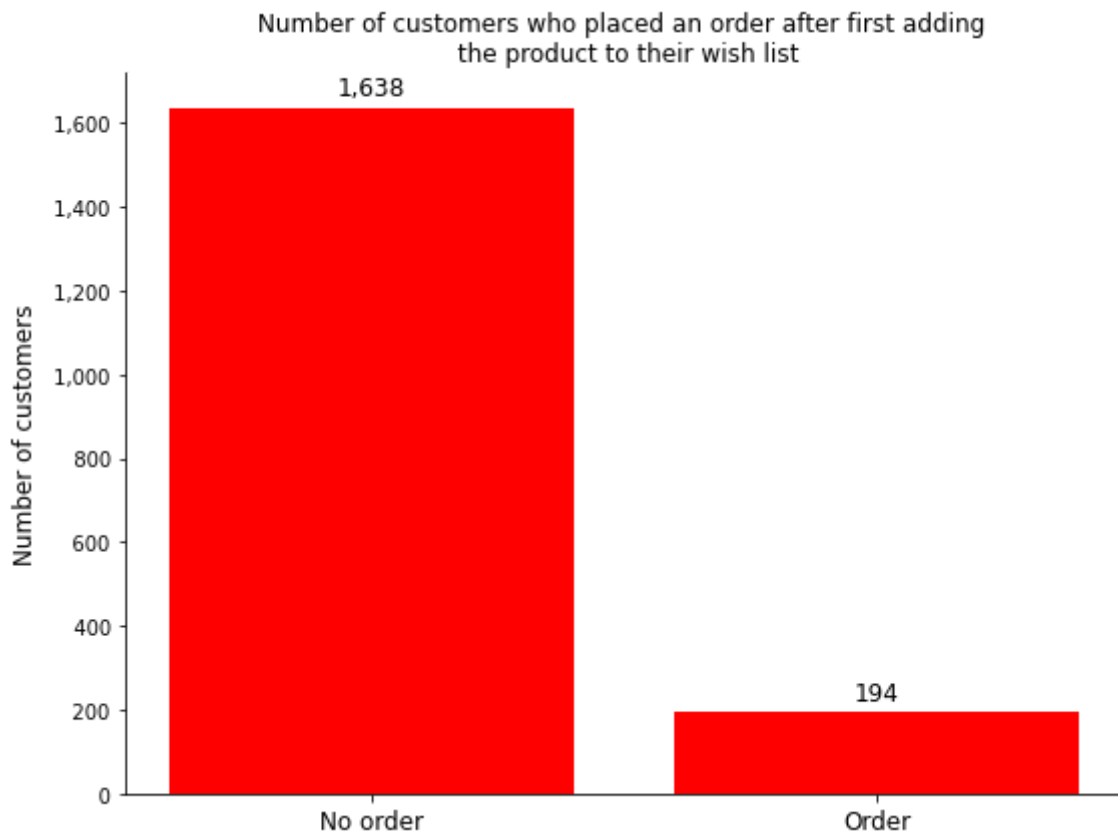
for bar in bars:
    yval = bar.get_height()
    #the '{:,}' command adds a thousand separator to the labels
    ax.annotate('{:,}'.format(yval),
        xy = (bar.get_x() + bar.get_width() / 2, yval),
        #shows label position on x and y axis
        xytext = (0, 3),
        textcoords = 'offset points', ha = 'center', va = 'bottom', fontsize = 12)

#function to add comma separator to labels
def comma(x, pos):
    return format(x, "6,.0f")

#this code adds a comma separator to the y tick marks
ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))

plt.xticks(wishlist['ordered'], labels = ['No order', 'Order'], fontsize = 12)
plt.ylabel('Number of customers', fontsize = 12)
plt.title('Number of customers who placed an order after first adding \n the product to
their wish list', y = 1.0, fontsize = 12)
plt.tight_layout()

plt.savefig('wishlist_volumes')
```



In [87]:

```
#grouping by customers who are returning customers
returning = data.groupby('returning_user').size().sort_values(ascending = False).to_frame().reset_index().\
rename(columns = {0: 'count'})
returning
```

Out[87]:

	returning_user	count
0	1	155441
1	0	76798

In [88]:

```
#grouping by customers who are returning customers and placed an order
returning_ordered = data.groupby(['returning_user', 'ordered']).size().to_frame().reset_index().\
rename(columns = {0: 'count'})
returning_ordered
```

Out[88]:

	returning_user	ordered	count
0	0	0	70663
1	0	1	6135
2	1	0	142483
3	1	1	12958

In [89]:

```
#filtering for returning users
```

```
returning_ordered = returning_ordered[returning_ordered['returning_user'] == 1].reset_index(drop = True)  
returning_ordered
```

Out[89]:

	returning_user	ordered	count
0	1	0	142483
1	1	1	12958

In [150]:

```
#creating chart
fig = plt.figure(figsize = (15, 6))
ax = plt.subplot(1, 2, 1)

_, _ , autotexts = plt.pie(returning['count'], colors = ['red', 'black'], radius = 1.5,
    autopct = ('%.0f%%'), \
        counterclock = False, startangle = -270)

plt.legend(labels = ['Returning customer', 'New customer'], loc = 'right', bbox_to_anchor = (1.7, 0.7), fontsize = 12)

#setting the colour of percentage labels to white
for autotext in autotexts:
    autotext.set_color('white')
    autotext.set_weight('bold')
    autotext.set_fontsize(12)

plt.title('Proportion of returning customers at Pedal Power', y = 1.2, fontsize = 12)

#second chart
ax = plt.subplot(1, 2, 2)

plt.bar(returning_ordered['ordered'], returning_ordered['count'], color = 'red')

#removing chart borders
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

#adding data labels for bars
bars = plt.bar(returning_ordered['ordered'], returning_ordered['count'], color = 'red')

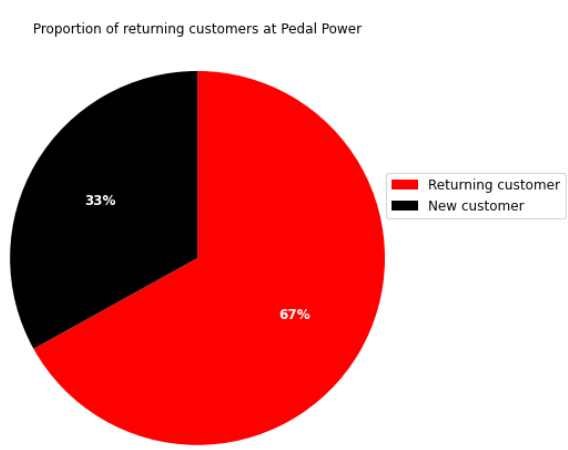
for bar in bars:
    yval = bar.get_height()
    #the '{:,}' command adds a thousand separator to the labels
    ax.annotate('{:,}'.format(yval),
        xy = (bar.get_x() + bar.get_width() / 2, yval),
        #shows label position on x and y axis
        xytext = (0, 3),
        textcoords = 'offset points', ha = 'center', va = 'bottom', fontsize = 12)

#function to add comma separator to labels
def comma(x, pos):
    return format(x, "6,.0f")

#this code adds a comma separator to the y tick marks
ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))

plt.xticks(returning_ordered['ordered'], labels = ['No order', 'Order'], fontsize = 12)
plt.ylabel('Number of customers', fontsize = 12)
plt.title('Number of returning customers who placed an order with \n Pedal Power', y = 1.2, fontsize = 12)
plt.tight_layout()

plt.savefig('returning_proportion_and_volumes')
```



In [91]:

```
#grouping by customers who went to the checkout page
checkout = data.groupby('saw_checkout').size().sort_values(ascending = False).to_frame()
.reset_index().\
rename(columns = {0: 'count'})
checkout
```

Out[91]:

	saw_checkout	count
0	0	191159
1	1	41080

In [92]:

```
#grouping by customers who saw the checkout page and placed an order
checkout_ordered = data.groupby(['saw_checkout', 'ordered']).size().to_frame().reset_in
dex().\
rename(columns = {0: 'count'})
checkout_ordered
```

Out[92]:

	saw_checkout	ordered	count
0	0	0	191159
1	1	0	21987
2	1	1	19093

In [93]:

```
#filtering for users who saw the checkout page
checkout_ordered = checkout_ordered[checkout_ordered['saw_checkout'] == 1].reset_index(
drop = True)
checkout_ordered
```

Out[93]:

	saw_checkout	ordered	count
0	1	0	21987
1	1	1	19093

In [151]:

```

#creating chart
fig = plt.figure(figsize = (15, 6))
ax = plt.subplot(1, 2, 1)

_, _ = plt.pie(checkout['count'], colors = ['red', 'black'], radius = 1.5, autopct = ('%.0f%%'), \
               counterclock = False, startangle = -270)

plt.legend(labels = ['Did not click checkout page', 'Clicked checkout page'], loc = 'right', bbox_to_anchor = (1.8, 1.0), \
           fontsize = 12)

#setting the colour of percentage labels to white
for autotext in autotexts:
    autotext.set_color('white')
    autotext.set_weight('bold')
    autotext.set_fontsize(12)

plt.title('Proportion of customers who viewed the checkout page \n at Pedal Power', y = 1.2, fontsize = 12)

#second chart
ax = plt.subplot(1, 2, 2)

plt.bar(checkout_ordered['ordered'], checkout_ordered['count'], color = 'red')

#removing chart borders
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

#adding data labels for bars
bars = plt.bar(checkout_ordered['ordered'], checkout_ordered['count'], color = 'red')

for bar in bars:
    yval = bar.get_height()
    #the '{:,}' command adds a thousand separator to the labels
    ax.annotate('{:,}'.format(yval),
                xy = (bar.get_x() + bar.get_width() / 2, yval),
                #shows label position on x and y axis
                xytext = (0, 3),
                textcoords = 'offset points', ha = 'center', va = 'bottom', fontsize = 12)

#function to add comma separator to labels
def comma(x, pos):
    return format(x, "6,.0f")


#this code adds a comma separator to the y tick marks
ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))

plt.xticks(checkout_ordered['ordered'], labels = ['No order', 'Order'], fontsize = 12)
plt.ylabel('Number of customers', fontsize = 12)
plt.title('Number of customers who viewed the checkout page and \n placed an order with Pedal Power', \
          y = 1.2, fontsize = 12)
plt.tight_layout()

plt.savefig('saw_checkout_proportion_and_volumes')

```

Proportion of customers who viewed the checkout page at Pedal Power



Number of customers who viewed the checkout page and placed an order with Pedal Power




Correlation analysis

In [95]:

```
#plot correlations between target feature (ordered) and all other features - checking
#delivery and viewing the checkout\
#page is most positively correlated with orders and list dropdown is most negatively co
rrelated
fig = plt.figure(figsize = (15, 8))
ax = plt.subplot()
data.corr()['ordered'].sort_values(ascending = False).plot(kind = 'bar', color = 'red')
plt.title('Feature correlation with orders placed', fontsize = 12)
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)
plt.xticks(fontsize = 10)
plt.yticks(fontsize = 10)
plt.tight_layout()
plt.savefig('correlation')
```

Feature correlation with orders placed



In [96]:

```
#splitting data into dependent and independent variables
y = data.ordered.values

X = data.drop('ordered', axis = 1)

#saving the X value columns to a separate list for reassigning after scaling the data
data_columns = X.columns
```

In [97]:

```
#splitting data into training and testing datasets - we are dealing with only 0/1 numbers so no need to scale
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
```

In [98]:

```
#checking shape of training and testing sets - all records are there
print('X_train shape:', X_train.shape)
print('X_test shape:', X_test.shape)
print('y_train shape:', y_train.shape)
print('y_test shape:', y_test.shape)
```

```
X_train shape: (185791, 19)
X_test shape: (46448, 19)
y_train shape: (185791,)
y_test shape: (46448,)
```

Logistic regression

In [99]:

```
#fitting the model - using a sag solver as this is faster for large datasets which we have
logistic_model = LogisticRegression(solver = 'sag', random_state = 42)
logistic_model.fit(X_train, y_train)
```

Out[99]:

```
LogisticRegression(random_state=42, solver='sag')
```

In [100]:

```
#predicting the purchase values for X_train and X_test
y_pred_train = logistic_model.predict(X_train)
y_pred_test = logistic_model.predict(X_test)
```

In [101]:

```
#evaluating the model - finding the residual differences between train data and predicted train data
residuals = np.abs(y_train - y_pred_train)

#print value counts of predicted values
print(pd.Series(residuals).value_counts())

print('')

#print normalised value counts for predicted values - 98% of order predictions were correct on the train data set
print(pd.Series(residuals).value_counts(normalize = True))
```

```
0    182591
1      3200
dtype: int64

0    0.982776
1    0.017224
dtype: float64
```

In [102]:

```
#evaluating the model for the test data - finding the residual differences between test data and predicted test data
residuals_test = np.abs(y_test - y_pred_test)

#print value counts of predicted values
print(pd.Series(residuals_test).value_counts())

print('')

#print normalised value counts for predicted values - 98% of order predictions were correct on the test data set
print(pd.Series(residuals_test).value_counts(normalize = True))
```

```
0    45705
1      743
dtype: int64

0    0.984004
1    0.015996
dtype: float64
```


In [103]:

```
#definitions
#precision - how precise the predictions are or  $TP / (TP + FP)$ . (Out of the times the model said the customer would order,
#how many times did they actually order)

#recall - what percentage of the class we're interested in were captured by the model or  $TP / (TP + FN)$ . (Out of all the
#customers that ordered, what percentage did the model predict as 'going to order')

#accuracy - measures what percentage of predictions the model got right or  $(TP + TN) / (TP + FP + TN + FN)$ .

#F1 score - harmonic mean of precision and recall - can't have a high F1 score without a strong model underneath.
#F1 =  $2(precision * recall) / (precision + recall)$ 
#F1 score penalises model heavily if it's skewed towards precision or recall
```

In [104]:

```
#printing the confusion matrix - we are trying to limit false positives (i.e. when we predict a customer will place an order
#but they don't). When trying to limit false positives we want to optimise on precision rather than recall

#41957 = true negative, 3748 = true positive, 56 = false negative, 687 = false positive
matrix = confusion_matrix(y_test, y_pred_test)
print('Confusion matrix \n', matrix)
```

Confusion matrix

```
[[41957  687]
 [   56 3748]]
```

In [105]:

```
#visualising the confusion matrix with labels
y_test_df = pd.Series(y_test)
y_pred_test_df = pd.Series(y_pred_test)

pd.crosstab(y_test_df, y_pred_test_df, rownames = ['True'], colnames = ['Predicted'], margins = True).reset_index()
```

Out[105]:

	Predicted	True	0	1	All
0	0	0	41957	687	42644
1	1	1	56	3748	3804
2	All	All	42013	4435	46448

In [106]:

```

#concatenating the series into a dataframe and adding a new row to describe what the re
sult is in each. The function uses
#if statements to define the result

results = pd.concat([y_test_df, y_pred_test_df], axis = 1).rename(columns = {0: 'actua
l', 1: 'prediction'})

def regression_results(row):
    if row['actual'] == 1 and row['prediction'] == 1:
        return 'actual = ordered and prediction = ordered - True positive'
    if row['actual'] == 0 and row['prediction'] == 0:
        return 'actual = did not order and prediction = did not order - True negative'
    if row['actual'] == 0 and row['prediction'] == 1:
        return 'actual = did not order but prediction = ordered - False positive'
    if row['actual'] == 1 and row['prediction'] == 0:
        return 'actual = ordered but prediction = did not order - False negative'

#applying the above function to a new row
results['result'] = results.apply(lambda row: regression_results(row), axis = 1)
print(results.head())

```

	actual	prediction	result
0	0	0	actual = did not order and prediction = did n...
1	0	0	actual = did not order and prediction = did n...
2	0	0	actual = did not order and prediction = did n...
3	0	0	actual = did not order and prediction = did n...
4	0	0	actual = did not order and prediction = did n...

In [107]:

```

#grouping the results to compare with the confusion matrix
results_grouped = results.groupby('result').size()
results_grouped

```

Out[107]:

```

result
actual = did not order and prediction = did not order - True negative
41957
actual = did not order but prediction = ordered - False positive
687
actual = ordered and prediction = ordered - True positive
3748
actual = ordered but prediction = did not order - False negative
56
dtype: int64

```

In [108]:

```
#printing out the scores for precision, recall, accuracy and F1 for y_test and y_pred_t
est
print(metrics.classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	42644
1	0.85	0.99	0.91	3804
accuracy			0.98	46448
macro avg	0.92	0.98	0.95	46448
weighted avg	0.99	0.98	0.98	46448

In [109]:

```
#another way to print the scores, pos_label tells sklearn what class you want to print
- we want 1 as that's the 'will
#order' class
```

```
precision_train = precision_score(y_train, y_pred_train, pos_label = 1)
precision_test = precision_score(y_test, y_pred_test, pos_label = 1)
```

```
recall_train = recall_score(y_train, y_pred_train, pos_label = 1)
recall_test = recall_score(y_test, y_pred_test, pos_label = 1)
```

```
accuracy_train = accuracy_score(y_train, y_pred_train)
accuracy_test = accuracy_score(y_test, y_pred_test)
```

```
f1_train = f1_score(y_train, y_pred_train, pos_label = 1)
f1_test = f1_score(y_test, y_pred_test, pos_label = 1)
```

```
print('Precision train: ', round(precision_train, 2))
print('Precision test: ', round(precision_test, 2))
print('')
print('Recall train: ', round(recall_train, 2))
print('Recall test: ', round(recall_test, 2))
print('')
print('Accuracy train: ', round(accuracy_train, 2))
print('Accuracy test: ', round(accuracy_test, 2))
print('')
print('F1 train: ', round(f1_train, 2))
print('F1 test: ', round(f1_test, 2))
print('')
```

```
Precision train: 0.84
Precision test: 0.85
```

```
Recall train: 0.98
Recall test: 0.99
```

```
Accuracy train: 0.98
Accuracy test: 0.98
```

```
F1 train: 0.9
F1 test: 0.91
```

In [110]:

```
#calculating probabilities scores for test and train sets
prob_train = logistic_model.predict_proba(X_train)
prob_test = logistic_model.predict_proba(X_test)

#keeping probabilities for positive outcome only (the threshold is 0.5 which means, if
the predicted probability of the
#class for an instance is less than 0.5, that instance is predicted to be in class 0 (t
he negative class). If the
#probability of the class for an instance is equal or greater than 0.5, the instance is
classified as class 1.)
prob_train_order = prob_train[:, 1]
prob_test_order = prob_test[:, 1]

#calculate false positive rate (fpr), true positive rate (tpr) and thresholds for train
set
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, prob_train_order)

#calculate false positive rate (fpr), true positive rate (tpr) and thresholds for test
set
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, prob_test_order)
```

In [111]:


```
#calculating AUC (Area Under Curve) - this gives a single numeric metric to evaluate th
e model. An AUC value of 1 would
#represent a perfect classifier. If AUC = 0.5 the classifier only has 50% preision. AUC
scores above 0.99 so pretty good.

auc_train = auc(train_fpr, train_tpr)
auc_test = auc(test_fpr, test_tpr)
print('Training AUC: ', round(auc_train, 5))
print('Testing AUC: ', round(auc_test, 5))
```

Training AUC: 0.99289
Testing AUC: 0.99328

In [112]:

```
#plotting an ROC (Receiver Operator Characteristic) curve to evaluate the model performance visually. It illustrates the  
#true positive rate against the false positive rate of our classifier. Best performing  
models will have an ROC curve that  
#hugs the upper left corner of the graph. This would represent that we correctly classify  
the positives much more often  
#than we incorrectly classify them.  
  
fig = plt.figure(figsize = (12, 6))  
ax = plt.subplot(1, 2, 1)  
plt.plot(train_fpr, train_tpr, label = 'ROC curve', color = 'red')  
  
#plotting diagonal line from zero which represents a model that guesses incorrectly for  
every correct guess  
plt.plot([0, 1], [0, 1], color = 'black', linestyle = 'dashed')  
plt.xlim([0, 1])  
plt.ylim([0, 1.05])  
  
plt.legend()  
plt.xticks(fontsize = 12)  
plt.yticks(fontsize = 12)  
plt.xlabel('False positive rate', fontsize = 12)  
plt.ylabel('True positive rate', fontsize = 12)  
plt.title('ROC curve for training set', fontsize = 12)  
  
ax = plt.subplot(1, 2, 2)  
plt.plot(test_fpr, test_tpr, label = 'ROC curve', color = 'red')  
  
#plotting diagonal line from zero which represents a model that guesses incorrectly for  
every correct guess  
plt.plot([0, 1], [0, 1], color = 'black', linestyle = 'dashed')  
plt.xlim([0, 1])  
plt.ylim([0, 1.05])  
  
plt.legend()  
plt.xticks(fontsize = 12)  
plt.yticks(fontsize = 12)  
plt.xlabel('False positive rate', fontsize = 12)  
plt.ylabel('True positive rate', fontsize = 12)  
plt.title('ROC curve for test set', fontsize = 12)  
plt.savefig('ROC_curves')
```



In [113]:

```
#the ROC curve shows a near perfect classifier. Looking at the data to double check this is correct - it is - these figures
#match the confusion matrix
print('counts for test set:')
print(pd.Series(y_test).value_counts())
print('')
print('counts for predictions on test set:')
print(pd.Series(y_pred_test).value_counts())
```

counts for test set:

0 42644

1 3804

dtype: int64

counts for predictions on test set:

0 42013

1 4435

dtype: int64

Random Forest

In [114]:

```
#first running the data through a random forest model with default hyperparameters
default_forest = RandomForestClassifier(random_state = 42)
default_model = default_forest.fit(X_train, y_train)
default_y_pred = default_model.predict(X_test)
```

In [115]:

```
#printing the confusion matrix - 669 false positives versus 687 with the logistic regression so random forest performed
#better in reducing false positives

#41975 = true negative, 3703 = true positive, 101 = false negative, 669 = false positive
matrix = confusion_matrix(y_test, default_y_pred)
print('Confusion matrix \n', matrix)
```

```
Confusion matrix
[[41975  669]
 [  101 3703]]
```

In [116]:

```
#printing out the scores for precision, recall, accuracy and F1 for y_test and default_y_pred - accuracy and precision
#are same as logistic regression
print(metrics.classification_report(y_test, default_y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	42644
1	0.85	0.97	0.91	3804
accuracy			0.98	46448
macro avg	0.92	0.98	0.95	46448
weighted avg	0.99	0.98	0.98	46448

In [117]:

```
#will tune some of the hyperparameters of the random forest model to see if scores can be improved. will tune number of
#estimators, max depth, min sample split and min sample leaf. Starting with n_estimators - this shows the number of trees in
#the forest. Default value is 10

num_est = [50, 100, 250, 500, 750]

#preparing to create a validation curve - this is a tool that shows how accuracy changes with changes in the model
#parameters.
#cv = cross validation - it's used to test the effectiveness of machine learning models and is a resampling procedure
train_score_num_est, test_score_num_est = validation_curve(
    RandomForestClassifier(),
    X = X_train, y = y_train,
    param_name = 'n_estimators',
    param_range = num_est,
    cv = 5,
    scoring = 'accuracy')
```

In [118]:

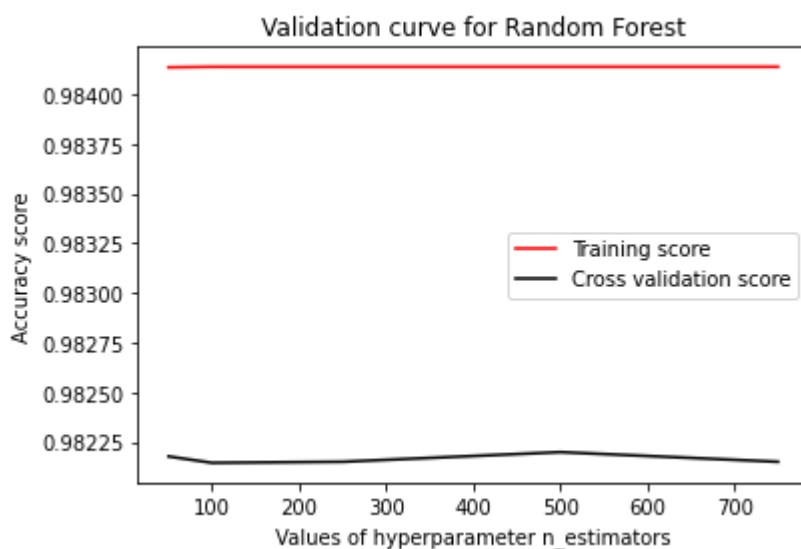
```
#to plot the validation curve, we need the mean of the 5 cross validation scores - calculating mean and standard deviation  
#for train and test sets  
  
mean_train_score_num_est = np.mean(train_score_num_est, axis = 1)  
std_train_score_num_est = np.std(train_score_num_est, axis = 1)  
  
mean_test_score_num_est = np.mean(test_score_num_est, axis = 1)  
std_test_score_num_est = np.std(test_score_num_est, axis = 1)
```

In [119]:

```
#plotting the validation curve - Large difference between training scores and cross validation scores but the training set  
#still had an average accuracy of 98% for each of the three cross validations. The two lines are ever so slightly closer  
#at 50 estimators so will use this in tuned model  
  
plt.plot(num_est, mean_train_score_num_est, label = 'Training score', color = 'red')  
plt.plot(num_est, mean_test_score_num_est, label = 'Cross validation score', color = 'black')  
  
plt.title('Validation curve for Random Forest')  
plt.xlabel('Values of hyperparameter n_estimators')  
plt.ylabel('Accuracy score')  
plt.legend(loc = 'best')
```

Out[119]:

<matplotlib.legend.Legend at 0x7fa1eaa3c358>



In [120]:

```
#tuning max_depth which specifies the maximum depth of each tree. The default is none w  
hich means that each tree will  
#expand until each leaf is pure, i.e. where all the data on the leaf comes from the sam  
e class
```

```
depth = [1, 5, 10, 20, 30]
```

```
#preparing to create the cross validation curve
```

```
train_score_depth, test_score_depth = validation_curve(  
    RandomForestClassifier(),  
    X = X_train, y = y_train,  
    param_name = 'max_depth',  
    param_range = depth,  
    cv = 5,  
    scoring = 'accuracy')
```

In [121]:

```
#calculating mean and standard deviation for train and test sets
```

```
mean_train_score_depth = np.mean(train_score_depth, axis = 1)  
std_train_score_depth = np.std(train_score_depth, axis = 1)
```

```
mean_test_score_depth = np.mean(test_score_depth, axis = 1)  
std_test_score_depth = np.std(test_score_depth, axis = 1)
```

In [122]:


```
#both the training and cross validations scores are highst at a tree depth of 5 so this  
will be used in the final model
```

```
plt.plot(depth, mean_train_score_depth, label = 'Training score', color = 'red')  
plt.plot(depth, mean_test_score_depth, label = 'Cross validation score', color = 'black')
```

```
plt.title('Validation curve for Random Forest')  
plt.xlabel('Values of hyperparameter max_depth')  
plt.ylabel('Accuracy score')  
plt.legend(loc = 'best')
```

Out[122]:

<matplotlib.legend.Legend at 0x7fa1eb56b7b8>



In [123]:

```
#tuning min_samples_split which is the minimum number of samples required to split an internal leaf node.  
#The default value is 2 which means that each internal leaf node must have at least two samples before it can split into a  
#more specific classification  
  
sample_split = [5, 10, 15, 20, 25]  
  
#preparing to create the cross validation curve  
  
train_score_split, test_score_split = validation_curve(  
    RandomForestClassifier(),  
    X = X_train, y = y_train,  
    param_name = 'min_samples_split',  
    param_range = sample_split,  
    cv = 5,  
    scoring = 'accuracy')
```

In [124]:


```
#calculating mean and standard deviation for train and test sets  
  
mean_train_score_split = np.mean(train_score_split, axis = 1)  
std_train_score_split = np.std(train_score_split, axis = 1)  
  
mean_test_score_split = np.mean(test_score_split, axis = 1)  
std_test_score_split = np.std(test_score_split, axis = 1)
```

In [125]:

```
#cross validation score is highest at 10 samples and then starts to flatten out. At this point accuracy is falling but it's  
#still high at 98.3%. Will therefore go with 10 samples before a leaf node can split - this should be ok given that we  
#have over 100,000 samples in the training set  
  
plt.plot(sample_split, mean_train_score_split, label = 'Training score', color = 'red')  
plt.plot(sample_split, mean_test_score_split, label = 'Cross validation score', color = 'black')  
  
plt.title('Validation curve for Random Forest')  
plt.xlabel('Values of hyperparameter min_samples_split')  
plt.ylabel('Accuracy score')  
plt.legend(loc = 'best')
```

Out[125]:

<matplotlib.legend.Legend at 0x7fa1eaecd668>



In [126]:

```
#tuning min_samples_leaf which is the minimum number of samples required to be a leaf node. The default is 1 which means  
#that every leaf must have at least one sample that it classifies  
  
sample_leaf = [2, 4, 6, 8, 10]  
  
#preparing to create the cross validation curve  
  
train_score_leaf, test_score_leaf = validation_curve(  
    RandomForestClassifier(),  
    X = X_train, y = y_train,  
    param_name = 'min_samples_leaf',  
    param_range = sample_leaf,  
    cv = 5,  
    scoring = 'accuracy')
```

In [127]:


```
#calculating mean and standard deviation for train and test sets  
  
mean_train_score_leaf = np.mean(train_score_leaf, axis = 1)  
std_train_score_leaf = np.std(train_score_leaf, axis = 1)  
  
mean_test_score_leaf = np.mean(test_score_leaf, axis = 1)  
std_test_score_leaf = np.std(test_score_leaf, axis = 1)
```

In [128]:

```
#cross validation score is highest at 6 samples. At this point accuracy is falling but  
it's still high at 98.3%. Will  
#therefore go with 8 samples as the minumum for each leaf node to classify  
  
plt.plot(sample_leaf, mean_train_score_leaf, label = 'Training score', color = 'red')  
plt.plot(sample_leaf, mean_test_score_leaf, label = 'Cross validation score', color =  
'black')  
  
plt.title('Validation curve for Random Forest')  
plt.xlabel('Values of hyperparameter min_samples_leaf')  
plt.ylabel('Accuracy score')  
plt.legend(loc = 'best')
```

Out[128]:

<matplotlib.legend.Legend at 0x7fa1eaa4de10>



In [139]:

```
#now putting our new hyperparameters into the final model

tuned_forest = RandomForestClassifier(n_estimators = 50, max_depth = 5, min_samples_split = 10, min_samples_leaf = 8,
                                     random_state = 42)

tuned_model = tuned_forest.fit(X_train, y_train)

tuned_y_pred = tuned_model.predict(X_test)
```

In [140]:

```
#printing the confusion matrix - 711 false positives versus 669 in default forest

#41933 = true negative, 3752 = true positive, 52 = false negative, 711 = false positive
matrix = confusion_matrix(y_test, tuned_y_pred)
print('Confusion matrix \n', matrix)
```

Confusion matrix

```
[[41933  711]
 [  52 3752]]
```

In [141]:

```
#printing out the scores for precision, recall, accuracy and F1 for y_test and default_y_pred - accuracy and precision
#are same as logistic regression and default random forest when printed to 2 dp
print(metrics.classification_report(y_test, tuned_y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	42644
1	0.84	0.99	0.91	3804
accuracy			0.98	46448
macro avg	0.92	0.98	0.95	46448
weighted avg	0.99	0.98	0.98	46448

In [142]:

```
#printing all scores for all models for class 1 on the test sets. Default model has a higher precision score than the tuned
#model and logistic regression model also had a higher precision score then tuned model. Logistic regression model also has
#a higher accuracy and F1 score than both random forest models although differences are small

#precision
log_reg_precision = precision_score(y_test, y_pred_test, pos_label = 1)
default_forest_precision = precision_score(y_test, default_y_pred, pos_label = 1)
tuned_forest_precision = precision_score(y_test, tuned_y_pred, pos_label = 1)

#recall
log_reg_recall = recall_score(y_test, y_pred_test, pos_label = 1)
default_forest_recall = recall_score(y_test, default_y_pred, pos_label = 1)
tuned_forest_recall = recall_score(y_test, tuned_y_pred, pos_label = 1)

#accuracy
log_reg_accuracy = accuracy_score(y_test, y_pred_test)
default_forest_accuracy = accuracy_score(y_test, default_y_pred)
tuned_forest_accuracy = accuracy_score(y_test, tuned_y_pred)

#f1 score
log_reg_f1 = f1_score(y_test, y_pred_test, pos_label = 1)
default_forest_f1 = f1_score(y_test, default_y_pred, pos_label = 1)
tuned_forest_f1 = f1_score(y_test, tuned_y_pred, pos_label = 1)

print('Logistic regression precision: ', round(log_reg_precision, 5))
print('Default Random Forest precision: ', round(default_forest_precision, 5))
print('Tuned Random Forest precision: ', round(tuned_forest_precision, 5))
print('')
print('Logistic regression recall: ', round(log_reg_recall, 5))
print('Default Random Forest recall: ', round(default_forest_recall, 5))
print('Tuned Random Forest recall: ', round(tuned_forest_recall, 5))
print('')
print('Logistic regression accuracy: ', round(log_reg_accuracy, 5))
print('Default Random Forest accuracy: ', round(default_forest_accuracy, 5))
print('Tuned Random Forest accuracy: ', round(tuned_forest_accuracy, 5))
print('')
print('Logistic regression F1: ', round(log_reg_f1, 5))
print('Default Random Forest F1: ', round(default_forest_f1, 5))
print('Tuned Random Forest F1: ', round(tuned_forest_f1, 5))
```

```
Logistic regression precision: 0.8451
Default Random Forest precision: 0.84698
Tuned Random Forest precision: 0.84069
```

```
Logistic regression recall: 0.98528
Default Random Forest recall: 0.97345
Tuned Random Forest recall: 0.98633
```

```
Logistic regression accuracy: 0.984
Default Random Forest accuracy: 0.98342
Tuned Random Forest accuracy: 0.98357
```

```
Logistic regression F1: 0.90982
Default Random Forest F1: 0.90582
Tuned Random Forest F1: 0.90771
```


In [143]:

```
#calculating probabilities scores for test sets
default_prob = default_model.predict_proba(X_test)
tuned_prob = tuned_model.predict_proba(X_test)

#keeping probabilities for positive outcome only (the threshold is 0.5 which means, if
the predicted probability of the
#class for an instance is less than 0.5, that instance is predicted to be in class 0 (t
he negative class). If the
#probability of the class for an instance is equal or greater than 0.5, the instance is
classified as class 1.)
default_prob_order = default_prob[:, 1]
tuned_prob_order = tuned_prob[:, 1]

#calculate false positive rate (fpr), true positive rate (tpr) and thresholds for test
set
test_fpr_d, test_tpr_d, test_thresholds_d = roc_curve(y_test, default_prob_order)
test_fpr_t, test_tpr_t, test_thresholds_t = roc_curve(y_test, tuned_prob_order)
```

In [144]:

```
#calculating AUC (Area Under Curve) - AUC is higher for the tuned model. But Logistic r
egression model is still higher than
#both (0.99328)

default_auc = auc(test_fpr_d, test_tpr_d)
tuned_auc = auc(test_fpr_t, test_tpr_t)
print('Default AUC: ', round(default_auc, 5))
print('Tuned_AUC: ', round(tuned_auc, 5))
```

Default AUC: 0.99187

Tuned_AUC: 0.9926

In [145]:

*#plotting an ROC curves for random forest models - again they show a very good model, s
imilar to the logistic regression*

```
fig = plt.figure(figsize = (12, 6))
ax = plt.subplot(1, 2, 1)
plt.plot(test_fpr_d, test_tpr_d, label = 'ROC curve', color = 'red')
```

*#plotting diagonal line from zero which represents a model that guesses incorrectly for
every correct guess*

```
plt.plot([0, 1], [0, 1], color = 'black', linestyle = 'dashed')
plt.xlim([0, 1])
plt.ylim([0, 1.05])
```

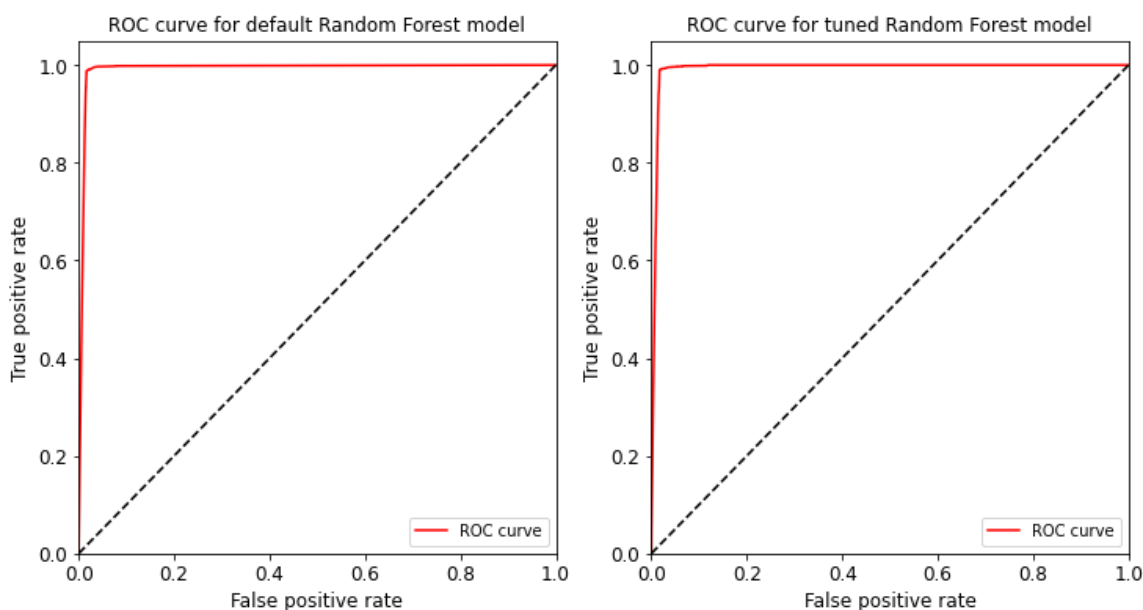
```
plt.legend()
plt.xticks(fontsize = 12)
plt.yticks(fontsize = 12)
plt.xlabel('False positive rate', fontsize = 12)
plt.ylabel('True positive rate', fontsize = 12)
plt.title('ROC curve for default Random Forest model', fontsize = 12)
```

```
ax = plt.subplot(1, 2, 2)
plt.plot(test_fpr_t, test_tpr_t, label = 'ROC curve', color = 'red')
```

*#plotting diagonal line from zero which represents a model that guesses incorrectly for
every correct guess*

```
plt.plot([0, 1], [0, 1], color = 'black', linestyle = 'dashed')
plt.xlim([0, 1])
plt.ylim([0, 1.05])
```

```
plt.legend()
plt.xticks(fontsize = 12)
plt.yticks(fontsize = 12)
plt.xlabel('False positive rate', fontsize = 12)
plt.ylabel('True positive rate', fontsize = 12)
plt.title('ROC curve for tuned Random Forest model', fontsize = 12)
plt.savefig('RF_ROC_curves')
```



In [146]:

#checking feature importance in default model

```
feature_importance = pd.Series(default_model.feature_importances_, index = X.columns)
feature_importance.sort_values(ascending = False, inplace = True)
print(feature_importance)
```

```
checked_delivery_detail    0.486037
saw_checkout               0.218194
sign_in                   0.152668
basket_add_detail          0.044707
basket_icon_click         0.042424
basket_add_list            0.012417
saw_delivery               0.007964
saw_homepage               0.006766
list_size_dropdown        0.005295
checked_returns_detail    0.005062
returning_user             0.004297
closed_minibasket_click   0.003183
image_picker              0.002450
promo_banner_click        0.002180
sort_by                   0.002065
detail_wishlist_add       0.001568
account_page_click        0.001540
saw_account_upgrade       0.000899
saw_sizecharts            0.000284
dtype: float64
```

In [147]:

#checking feature importance in tuned model

```
feature_importance_t = pd.Series(tuned_model.feature_importances_, index = X.columns)
feature_importance_t.sort_values(ascending = False, inplace = True)
print(feature_importance_t)
```

```
checked_delivery_detail    0.410433
saw_checkout               0.250969
sign_in                   0.200023
basket_add_detail          0.059542
basket_icon_click         0.044216
basket_add_list            0.013162
saw_homepage               0.006571
list_size_dropdown        0.004283
saw_delivery               0.003197
checked_returns_detail    0.002900
returning_user             0.001676
image_picker              0.001168
closed_minibasket_click   0.000848
detail_wishlist_add       0.000462
account_page_click        0.000291
sort_by                   0.000162
promo_banner_click        0.000044
saw_account_upgrade       0.000037
saw_sizecharts            0.000016
dtype: float64
```

In [148]:

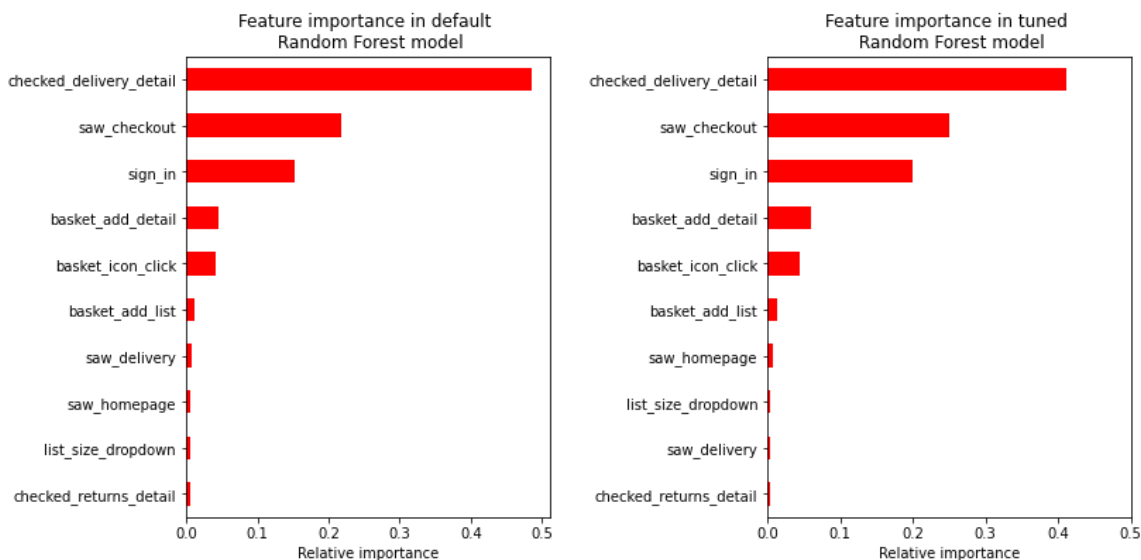
```
#plotting feature importance for default model
plt.figure(figsize = (12,6))
ax = plt.subplot(1, 2, 1)

feature_importance.nlargest(10).plot(kind = 'barh', color = 'red').invert_yaxis()
plt.xlabel('Relative importance')
plt.title('Feature importance in default \n Random Forest model', fontsize = 12)

ax = plt.subplot(1, 2, 2)

feature_importance_t.nlargest(10).plot(kind = 'barh', color = 'red').invert_yaxis()
plt.xlabel('Relative importance')
ax.set_xlim(0.0, 0.5)
plt.title('Feature importance in tuned \n Random Forest model', fontsize = 12)

plt.subplots_adjust(wspace = 0.6)
plt.savefig('feature_importance')
plt.show()
```



In []:

In []: